


```
In [ ]: #Name: Sushrut Deshpande  
#RBT23CB002  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import MinMaxScaler, LabelEncoder  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression, LinearRegression  
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [2]: raw_data = {  
    'Region': ['USA', 'China', 'Russia', 'India', 'UK', 'Australia', 'Europe', 'Brazil',  
    'Product': ['A', 'V', 'A', 'C', 'B', 'D', 'D', 'E', 'A'],  
    'Sales': [100, 125, 130, 200, 124, 245, 524, 10000, np.nan],  
    'Profit': [23, 42, 12, 42, 53, 123, 52, 5000, 23, np.nan]  
}
```


```
In [4]: df_1 = pd.DataFrame(raw_data)  
print(df_1)
```

	Region	Product	Sales	Profit
0	USA	A	100.0	23.0
1	China	V	125.0	42.0
2	Russia	A	130.0	12.0
3	India	C	200.0	42.0
4	UK	B	124.0	53.0
5	Australia	D	245.0	123.0
6	Europe	D	524.0	52.0
7	Brazil	D	10000.0	5000.0
8	South Africa	E	231.0	23.0
9	NaN	A	NaN	NaN

```
In [5]: df_1.plot(x='Region', y=['Sales', 'Profit'], kind='bar')  
plt.title('Sales and Profit by Region')  
plt.ylabel('Value')  
plt.xticks(rotation=45)  
plt.tight_layout()  
plt.show()
```



```
In [6]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.histplot(df_1['Sales'], bins=10, kde=True, ax=axes[0])
axes[0].set_title('Raw Data - Sales')
sns.histplot(df_1['Profit'], bins=5, kde=False, ax=axes[1])
axes[1].set_title('Raw Data - Profit')
plt.tight_layout()
plt.show()
df_1.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   Region    9 non-null      object  
 1   Product   10 non-null     object  
 2   Sales     9 non-null      float64 
 3   Profit    9 non-null      float64 
dtypes: float64(2), object(2)
memory usage: 452.0+ bytes
```

```
In [7]: df = df_1.copy()
df['Region'].fillna(df['Region'].mode()[0], inplace=True)
df['Profit'].fillna(df['Profit'].mean(), inplace=True)
df['Sales'].fillna(df['Sales'].mean(), inplace=True)
print(df)
```

	Region	Product	Sales	Profit
0	USA	A	100.000000	23.000000
1	China	V	125.000000	42.000000
2	Russia	A	130.000000	12.000000
3	India	C	200.000000	42.000000
4	UK	B	124.000000	53.000000
5	Australia	D	245.000000	123.000000
6	Europe	D	524.000000	52.000000
7	Brazil	D	10000.000000	5000.000000
8	South Africa	E	231.000000	23.000000
9	Australia	A	1297.666667	596.666667

C:\Users\USER\AppData\Local\Temp\ipykernel_14772\3058838888.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df['Region'].fillna(df['Region'].mode()[0], inplace=True)
C:\Users\USER\AppData\Local\Temp\ipykernel_14772\3058838888.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df['Profit'].fillna(df['Profit'].mean(), inplace=True)
C:\Users\USER\AppData\Local\Temp\ipykernel_14772\3058838888.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

df['Sales'].fillna(df['Sales'].mean(), inplace=True)

```
In [8]: for col in ['Sales', 'Profit']:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    print(Q1, Q3)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    df[col] = np.where(df[col] > upper, upper, np.where(df[col] < lower, lower, df[col]))
print(df)
```

```
126.25 454.25
27.75 105.5
      Region Product    Sales   Profit
0        USA       A  100.00  23.000
1     China       V  125.00  42.000
2    Russia       A  130.00  12.000
3     India       C  200.00  42.000
4       UK       B  124.00  53.000
5  Australia       D  245.00 123.000
6    Europe       D  524.00  52.000
7    Brazil       D  946.25 222.125
8  South Africa     E  231.00  23.000
9  Australia       A  946.25 222.125
```

```
In [9]: scaler = MinMaxScaler()
df[['Profit', 'Sales']] = scaler.fit_transform(df[['Profit', 'Sales']])
print(df)
```

```
      Region Product    Sales   Profit
0        USA       A  0.000000  0.052350
1     China       V  0.029542  0.142772
2    Russia       A  0.035451  0.000000
3     India       C  0.118168  0.142772
4       UK       B  0.028360  0.195122
5  Australia       D  0.171344  0.528257
6    Europe       D  0.501034  0.190363
7    Brazil       D  1.000000  1.000000
8  South Africa     E  0.154801  0.052350
9  Australia       A  1.000000  1.000000
```

```
In [10]: le_region = LabelEncoder()
le_product = LabelEncoder()


df['Region'] = le_region.fit_transform(df['Region'])
df['Product'] = le_product.fit_transform(df['Product'])

print(df)
```

```
      Region Product    Sales   Profit
0        8       0  0.000000  0.052350
1        2       5  0.029542  0.142772
2        5       0  0.035451  0.000000
3        4       2  0.118168  0.142772
4        7       1  0.028360  0.195122
5        0       3  0.171344  0.528257
6        3       3  0.501034  0.190363
7        1       3  1.000000  1.000000
8        6       4  0.154801  0.052350
9        0       0  1.000000  1.000000
```

```
In [11]: df.plot(x='Region', y=['Sales', 'Profit'], kind='bar')
plt.title('Sales and Profit by Region')
plt.ylabel('Value')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

Sales and Profit by Region



```
In [12]: fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.histplot(df['Sales'], bins=10, kde=True, ax=axes[0])
axes[0].set_title('Raw Data - Sales')
sns.histplot(df['Profit'], bins=10, kde=True, ax=axes[1])
axes[1].set_title('Raw Data - Profit')
plt.tight_layout()
plt.show()
```

